2020-2021 1-GS Methods in AI research (INFOMAIR)
Announcements
Dashboard
Staff Information
Course Information
Course Schedule
Course Content (lecture slides etc.)
Team Project
Questionnaire and submitting project deliverables
My Grades

Review Test Submission: Quiz Lecture 9 (Designing responsible AI) (only part by Dong)

User	Otto Mättas
Course	2020-2021 1-GS Methods in AI research (INFOMAIR)
Test	Quiz Lecture 9 (Designing responsible AI) (only part by Dong)
Started	9/29/20 2:56 PM
Submitted	9/29/20 3:05 PM
Due Date	9/30/20 4:00 PM
Status	Needs Grading
Attempt Score	Grade not available.
Time Elapsed	9 minutes
Results Displayed	All Answers, Submitted Answers, Correct Answers, Feedback, Incorrectly Answered Questions

Question 1

Needs Grading

In NLP there has been research on automatically estimating the demographics (e.g. gender, age) of people based on their language use, for example in social media.

Give an example of a potential 1) beneficial use and 2) harmful use.

Selected Answer:	1) Research to explain how diseases evolves 2) Models make decision-making general so individuals might fall under undeserved bias
Correct Answer:	There are many possible answers, but an example could be: 1) Beneficial use: To help social science research. For example, research may be interested in studying behavior in online platforms, but often researchers need to take demographics into account. Another example is analyzing public opinion, where knowing the demographics can help you get better estimates for the whole population. 2) Harmful use: Using this information for political influence (e.g. targeted messages). As another note: an additional risk is that misidentification (e.g. gender) can be quite harmful for individuals.
Response Feedback:	[None Given]

Question 2

Needs Grading

When someone says that a ML system is like Clever Hans, what do they mean with this?

Selected Answer:	We might reach a certain correct answer not necessarily using right signals. In seeking more answers, the invalid signals might fail to give an (correct) answer as they do not possess the capacity for it. They were simply not meant for the task at hand.
Correct Answer:	That the system hasn't actually learned the intended task even though performance numbers may be high. Instead, it has learned to use some shortcut/spurious signal.
Response Feedback:	[None Given]

Question 3

Needs Grading

Are any aspects of the lecture material unclear, do you have follow-up questions about this, or are there any topics you would like to learn more about?
I will try to address this by adding more background material to blackboard.

Leave this blank if you do not have any questions.

Selected Answer:	Thanks!
Correct Answer:	[None]
Response Feedback:	[None Given]

Wednesday, November 4, 2020 10:34:14 PM CET

← OK