

# From data to vectors & NLP 101

## Methods in AI research

Dong Nguyen  
17 Sept 2019



Utrecht University

# Practicalities

## Literature for today:

- Hal Daumé III, A Course in Machine Learning, 3.1-3.3 (Geometry and Nearest Neighbors; [http://ciml.info/dl/vo\\_99/ciml-vo\\_99-cho3.pdf](http://ciml.info/dl/vo_99/ciml-vo_99-cho3.pdf))
  - Description in text for Figure 3.4 is wrong (+ and – are switched)
- Jurafsky & Martin SLP3, 6.3 (Words and vectors) and 6.4 (Cosine for measuring similarity) <https://web.stanford.edu/~jurafsky/slp3/6.pdf>
- Noah A. Smith (2019), Contextual Word Representations: A Contextual Introduction: <https://arxiv.org/pdf/1902.06006.pdf> (sec. 5 optional).

# So far

- **ML concepts:**
  - Supervised learning (how to frame your task as a supervised learning problem)
  - Inductive bias
  - Overfitting and underfitting
  - Decision boundaries
  - Evaluation of supervised learning systems (don't touch your test data!)
- **Methods**
  - Decision Trees

# Today

- **ML concepts:**
  - Vector spaces
  - Distance metrics
- **ML method:**
  - K-nearest neighbors
- **NLP 101**
  - How to represent documents as vectors
  - How to represent words as vectors

## RECAP !

# Supervised learning

Learn a machine learning model using **labeled example instances**:

features                  target  
    ↘                  ↘  
     $\{<\mathbf{x}^{(1)}, \mathbf{y}^{(1)}>, \dots, <\mathbf{x}^{(N)}, \mathbf{y}^{(N)}>\}$

**Goal:** Predict the target using the features

Need to define **features**, characteristics of the instances that the model uses for predictions (words in a document, movie ratings, etc..)

## Features for house price prediction:


- Neighborhood
- Number of bedrooms
- First floor square meters
- Number of schools within 2 km
- Police Label Safe Housing
- ..

# Representing instances using feature vectors

Number of bedrooms: 2

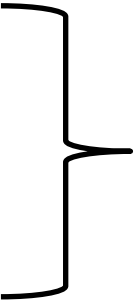
Plot size (square meters): 110

Number of schools within 2 km: 2



# Representing instances using feature vectors

Number of bedrooms: 2  
Plot size (square meters): 110  
Number of schools within 2 km: 2  
Police Label Safe Housing: yes



[ 2, 110, 2, 1 ]

*vector*

*We encode binary features as 1 (yes) and 0 (no)*

# Representing instances using feature vectors

Number of bedrooms: 2  
Plot size (square meters): 110  
Number of schools within 2 km: 2  
Police Label Safe Housing: yes  
**Neighborhood: Leidsche Rijn**

} [ 2, 110, 2, 1, ? ] *vector*

East: 0  
West: 1  
Leidsche Rijn: 2  
Overvecht: 3





# Representing instances using feature vectors

Number of bedrooms: 2  
Plot size (square meters): 110  
Number of schools within 2 km: 2  
Police Label Safe Housing: yes  
**Neighborhood: Leidsche Rijn**

} [ 2, 110, 2, 1, 0, 0, 1, 0 ]  
vector

East: 0  
West: 1  
Leidsche Rijn: 2  
Overvecht: 3



East? Yes = 1, No = 0    **One Hot Encoding**  
West? Yes = 1, No = 0  
Leidsche Rijn? Yes = 1, No = 0  
Overvecht: Yes = 1, No = 0




Neighborhood feature: [0,0,1,0]

# Types of features

- A **numerical** feature (a real number)
  - Sentence length
  - Number of likes
  - Temperature
- A **binary** feature: a yes vs. no distinction (usually 1 vs. 0)
  - Is the text capitalized?
  - Are A and B friends?
  - Employed?
  - Like Chinese restaurants?
- **Categorical** feature:
  - Country
  - Genre of a text

# Vector space



$$\mathbf{a} = [5, 5]$$

$$\mathbf{b} = [2, 1]$$

Your instances are now represented as points in *vector space*. Each dimension represents a feature (e.g. whether the user liked a certain restaurant)

Usually *thousands of features*

# Vectors & vector spaces

$$\mathbf{a} = [5, 5]$$

$$\mathbf{b} = [2, 1]$$

$\mathbf{a}$  is a two-dimensional vector, i.e.  $\mathbf{a} \in \mathbb{R}^2$


$$\mathbf{a} + \mathbf{b} = [5 + 2, 5 + 1] = [7, 6]$$

$$\mathbf{c} = [c_1, \dots, c_d]$$


$\mathbf{c}$  is a  $d$ -dimensional vector, i.e.  $\mathbf{c} \in \mathbb{R}^d$

What is  $a_1$ ? 5

What is  $b_2$ ? 1



# Vector space



**Q: How would you classify this point? (red or blue?)**

# Nearest neighbor

*This “**rule of nearest neighbor**” has considerable elementary intuitive appeal and probably corresponds to practice in many situations. For example, it is possible that much medical diagnosis is influenced by the doctor’s **recollection** of the subsequent history of an earlier patient whose symptoms **resemble in some way** those of the current patient. (Fix and Hodges, 1952)*

**Idea:** Classify new examples based on the most similar training examples

[CIML: chapter 3]

# Memory-based learner

This is **memory-based** learning (also called instance-based learning): look for similar instances in the training data (stored in **memory**) and fit with the local points.

## Four components:

- A distance metric
- How many neighbors to look at?
- A weighting function (*optional*)
- How to fit with the local points?

# Memory-based learner


This is **memory-based** learning (also called instance-based learning): look for similar instances in the training data (stored in **memory**) and fit with the local points.

**Four components:** 1-nearest neighbors

- A distance metric  
Many options, e.g. Euclidian
- How many neighbors to look at?  
1
- A weighting function (*optional*)  
Unused
- How to fit with the local points?  
Just return the label  
of the most nearest point




# 1-nearest neighbors decision boundaries



*Every training example has its own neighborhood*

For any point  $x$  in a training set, the Voronoi cell of  $x$  consists of all points closer to  $x$  than any other points in the training set.

# 1-nearest neighbors decision boundaries



1-nearest neighbors is sensitive to outliers!

Small changes in the training set can lead to large differences in the decision boundary

This point now gets classified as blue, is this what we want?

# Memory-based learner


This is **memory-based** learning (also called instance-based learning): look for similar instances in the training data (stored in **memory**) and fit with the local points.

## Four components: *K-nearest neighbors*

- A distance metric  
*Many options, e.g. Euclidian*
- How many neighbors to look at?  
*K*
- A weighting function (*optional*)  
*Unused*
- How to fit with the local points?  
*Just predict the majority label*

# K-nearest neighbor

Training data      number of neighbors      test instance



---

**Algorithm 3** KNN-PREDICT( $\mathbf{D}, K, \hat{\mathbf{x}}$ )

---

```
1:  $S \leftarrow []$ 
2: for  $n = 1$  to  $N$  do
3:    $S \leftarrow S \oplus \langle d(\mathbf{x}_n, \hat{\mathbf{x}}), n \rangle$            // store distance to training example  $n$ 
4: end for
5:  $S \leftarrow \text{SORT}(S)$                            // put lowest-distance objects first
6:  $\hat{y} \leftarrow 0$ 
7: for  $k = 1$  to  $K$  do
8:    $\langle \text{dist}, n \rangle \leftarrow S_k$                  //  $n$  this is the  $k$ th closest data point
9:    $\hat{y} \leftarrow \hat{y} + y_n$                        // vote according to the label for the  $n$ th training point
10: end for
11: return  $\text{SIGN}(\hat{y})$                              // return  $+1$  if  $\hat{y} > 0$  and  $-1$  if  $\hat{y} < 0$ 
```

---

[CIML: chapter 3]

# Choosing $K$

How many instances should we consider when making the classification?

$K = 1$

*Sensitive to outliers*

*Could lead to overfitting*


$K = N$  (number of instances)

*Always predict the majority label*


*Leads to underfitting!*

# Decision boundary

$K=1$

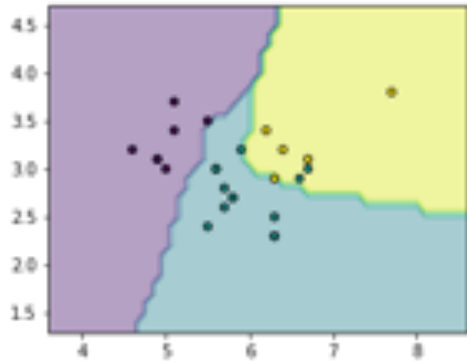


$K=15$



*Larger  $K$  values  
lead to smoother  
decision boundaries*



# Decision boundary



## Questions:

- Which one is:  $K=1$ ,  $K=3$ ,  $K=5$ ?
- $K$  is usually an odd number, why?
- How would you set  $K$ ?

# Decision boundary



## Questions:

- Which one is:  $K=1$ ,  $K=3$ ,  $K=5$ ?
- $K$  is usually an odd number, why?
- How would you set  $K$ ?

$K=5$ ,  $K=3$ ,  $K=1$

To not have ties (with binary classification)

$K$  is a hyperparameter



# K-nearest neighbor

Training data      number of neighbors      test instance



---

**Algorithm 3** KNN-PREDICT( $D, K, \hat{x}$ )

---

```
1:  $S \leftarrow []$ 
2: for  $n = 1$  to  $N$  do
3:    $S \leftarrow S \oplus \langle d(x_n, \hat{x}), n \rangle$ 
4: end for
5:  $S \leftarrow \text{SORT}(S)$ 
6:  $\hat{y} \leftarrow 0$ 
7: for  $k = 1$  to  $K$  do
8:    $\langle \text{dist}, n \rangle \leftarrow S_k$ 
9:    $\hat{y} \leftarrow \hat{y} + y_n$ 
10: end for
11: return  $\text{SIGN}(\hat{y})$ 
```

*// store distance to training example  $n$*

*// put lowest-distance objects first*

*//  $n$  this is the  $k$ th closest data point*

*// vote according to the label for the  $n$ th training point*

*// return  $+1$  if  $\hat{y} > 0$  and  $-1$  if  $\hat{y} < 0$*

How do we compute the distance between examples?

# Distance measure

By representing data points as vectors, we can measure their distance (or similarity) in the vector space.

If  $a$  and  $b$  are scalars, the most straightforward way to define distance is:

$$| a - b | \qquad (e.g., |3-5| = 2)$$

Let's generalize this idea to vectors

# Length (or, norm) of a vector

Length of vector  $[3, 4]$ :


$$\sqrt{x^2 + y^2} = \sqrt{4^2 + 3^2}$$

Length of vector  $[3, 4, 1]$ :

$$\sqrt{4^2 + 3^2 + 1^2}$$

Length of vector  $\mathbf{a}$ :

$$\|\mathbf{a}\|_2 = \sqrt{\sum a_i^2}$$



This is also called the  
 $l_2$ - norm or the  
Euclidian norm

# Euclidian distance

Distance between  
endpoint of vectors

Euclidian distance:

$$\|a - b\|_2 = \sqrt{\sum (a_i - b_i)^2}$$

## Question:

What is the Euclidian  
distance between a & b

a = [ 0 , 3 , 5 ]

b = [ 3 , 1 , 2 ]

# Euclidian distance

Distance between  
endpoint of vectors

$$\begin{aligned} a-b &= [(0-3), (3-1), (5-2)] \\ &= [-3, 2, 3] \end{aligned}$$

Euclidian distance:

$$\|a - b\|_2 = \sqrt{\sum (a_i - b_i)^2}$$

also  
called L2  
distance

$$\begin{aligned} &= \text{sqrt}((-3)^2 + 2^2 + 3^2) \\ &= \text{sqrt}(9 + 4 + 9) \\ &= 4,690 \end{aligned}$$

## Question:

What is the Euclidian  
distance between a & b

$$\begin{aligned} a &= [0, 3, 5] \\ b &= [3, 1, 2] \end{aligned}$$

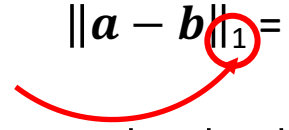
# Manhattan distance

Distance between  
endpoint of vectors

$$\begin{aligned} a-b &= [(0-3), (3-1), (5-2)] \\ &= [-3, 2, 3] \end{aligned}$$

Manhattan distance:

also  
called L1  
distance

$$\begin{aligned} \|a - b\|_1 &= \sum |a_i - b_i| \\ &= |-3| + |2| + |3| \\ &= 8 \end{aligned}$$



**Question:**

What is the Manhattan  
distance between a & b

$$a = [0, 3, 5]$$

$$b = [3, 1, 2]$$

# Comparison distances




**Minkowski distance:**  
Generalization of the Euclidian and  
Manhattan distance

$$\sqrt[p]{\sum |a_i - b_i|^p}$$

Many options!


Different distance measures →  
different neighborhoods



# Inductive bias of k-nearest neighbors

- The label of a data point should be similar to labels of nearby points
- All features are equally important (but weighted variants exist!)





# Feature Scaling

## Centering

the  $i$ th training instance

$$x_j^{(i)} = x_j^{(i)} - \mu_j$$

the  $j$ th feature

with  $\mu_j = \frac{1}{N} \sum_i x_j^{(i)}$

N training examples  
 $1 \leq i \leq N$   
 d features  
 $1 \leq j \leq d$

## Variance Scaling

$$x_j^{(i)} = \frac{x_j^{(i)}}{\sigma_d}$$

# Practical considerations

- How can we use K-Nearest Neighbors for **regression**?
  - Just predict the mean of the neighbors
- What about **ties**?
  - Random
  - Use 1 -nearest neighbor to decide
  - Use the class that is most frequent in the training data (highest prior probability)

# Computational Considerations

Training data      number of neighbors      test instance



---

**Algorithm 3** KNN-PREDICT( $\mathbf{D}, K, \hat{\mathbf{x}}$ )

---

```
1:  $S \leftarrow []$ 
2: for  $n = 1$  to  $N$  do
3:    $S \leftarrow S \oplus \langle d(\mathbf{x}_n, \hat{\mathbf{x}}), n \rangle$ 
4: end for
5:  $S \leftarrow \text{SORT}(S)$ 
6:  $\hat{y} \leftarrow 0$ 
7: for  $k = 1$  to  $K$  do
8:    $\langle \text{dist}, n \rangle \leftarrow S_k$ 
9:    $\hat{y} \leftarrow \hat{y} + y_n$ 
10: end for
11: return  $\text{SIGN}(\hat{y})$ 
```

Q: How does the classification speed (for test instances) depend on the **number of classes** in the training data?

// put lowest-distance objects first

//  $n$  this is the  $k$ th closest data point

// vote according to the label for the  $n$ th training point

// return  $+1$  if  $\hat{y} > 0$  and  $-1$  if  $\hat{y} < 0$

---

[CIML: chapter 3]

# Computational Considerations

Training data      number of neighbors      test instance



---

## Algorithm 3 KNN-PREDICT( $\mathbf{D}$ , $K$ , $\hat{x}$ )

---

```
1:  $S \leftarrow []$ 
2: for  $n = 1$  to  $N$  do
3:    $S \leftarrow S \oplus \langle d(\mathbf{x}_n, \hat{x}), n \rangle$ 
4: end for
5:  $S \leftarrow \text{SORT}(S)$ 
6:  $\hat{y} \leftarrow 0$ 
7: for  $k = 1$  to  $K$  do
8:    $\langle \text{dist}, n \rangle \leftarrow S_k$ 
9:    $\hat{y} \leftarrow \hat{y} + y_n$ 
10: end for
11: return  $\text{SIGN}(\hat{y})$ 
```

Q: How does the classification speed (for test instances) depend on the **number of classes** in the training data?  $\rightarrow$  *independent*

// put lowest-distance objects first


//  $n$  this is the  $k$ th closest data point

// vote according to the label for the  $n$ th training point

// return  $+1$  if  $\hat{y} > 0$  and  $-1$  if  $\hat{y} < 0$

# Computational Considerations

Training data      number of neighbors      test instance



---

**Algorithm 3** KNN-PREDICT( $\mathbf{D}, K, \hat{\mathbf{x}}$ )

---

```
1:  $S \leftarrow []$ 
2: for  $n = 1$  to  $N$  do
3:    $S \leftarrow S \oplus \langle d(\mathbf{x}_n, \hat{\mathbf{x}}), n \rangle$ 
4: end for
5:  $S \leftarrow \text{SORT}(S)$ 
6:  $\hat{y} \leftarrow 0$ 
7: for  $k = 1$  to  $K$  do
8:    $\langle \text{dist}, n \rangle \leftarrow S_k$ 
9:    $\hat{y} \leftarrow \hat{y} + y_n$ 
10: end for
11: return  $\text{SIGN}(\hat{y})$ 
```

Q: How does the classification speed (for test instances) depend on the **number of instances** in the training data?

// put lowest-distance objects first

//  $n$  this is the  $k$ th closest data point

// vote according to the label for the  $n$ th training point


// return  $+1$  if  $\hat{y} > 0$  and  $-1$  if  $\hat{y} < 0$

---

[CIML: chapter 3]

# Computational Considerations

Training data      number of neighbors      test instance



---

**Algorithm 3** KNN-PREDICT( $\mathbf{D}, K, \hat{\mathbf{x}}$ )

---

```
1:  $S \leftarrow []$ 
2: for  $n = 1$  to  $N$  do
3:    $S \leftarrow S \oplus \langle d(\mathbf{x}_n, \hat{\mathbf{x}}), n \rangle$ 
4: end for
5:  $S \leftarrow \text{SORT}(S)$ 
6:  $\hat{y} \leftarrow 0$ 
7: for  $k = 1$  to  $K$  do
8:    $\langle \text{dist}, n \rangle \leftarrow S_k$ 
9:    $\hat{y} \leftarrow \hat{y} + y_n$ 
10: end for
11: return  $\text{SIGN}(\hat{y})$ 
```

Q: How does the classification speed (for test instances) depend on the **number of instances** in the training data?  $\rightarrow$  linearly ☹️

// put lowest-distance objects first

//  $n$  this is the  $k$ th closest data point

// vote according to the label for the  $n$ th training point

// return  $+1$  if  $\hat{y} > 0$  and  $-1$  if  $\hat{y} < 0$

---

[CIML: chapter 3]

# Computational Considerations

- Training is fast
- It's easy to add new training data (no 'retraining' needed)



- Making predictions is slow
  - Takes  $N$  comparisons (number of instances)  $\times$   $d$  (number of dimensions) operations
- Need to *store* the training data



*There are techniques to speed up  $K$ -nearest neighbors, such as  $kd$ -trees*

# Decision trees vs. nearest neighbors

## Decision Trees

- Learn a model from the training data
- Apply model to new data
- Features are selected
- Decision boundaries are axis-aligned cuts

## Nearest neighbors

- Store data
- Compare new data to stored data
- All features have equal weight → Sensitive to irrelevant features
- Decision boundaries can be complex



# Chatbots Architectures

- **Rule-based**
  - Pattern-action rules (ELIZA) + a mental model (Parry)
- **Corpus-based (using a large chat corpus)**
  - Information Retrieval (IR) based. **Select** a response from a large corpus.
  - Neural networks. Train a neural network to generate text.

# Chatbots, response selection

## Return the response to the most similar turn

- Take user's turn **q** and find a similar turn **t** in the corpus **C**

**q** = "do you like Doctor Who?"

**t** = "you like Doctor Who?"

- Return the response to **t** in your corpus.

$$r = \text{response}(\operatorname{argmax}_{t \in C} \text{similarity}(q, t))$$

1-NN approach!

Up next: How can we represent words  
and documents as vectors?

# Natural Language Processing 101

# Natural Language Processing (NLP)

Automatic processing and analysis of  
**natural language**

Dutch, English, Spanish, Hindi, Frisian, Turkish, ...

7,111 known living languages.


<https://www.ethnologue.com/>

# Analysis of linguistic structure

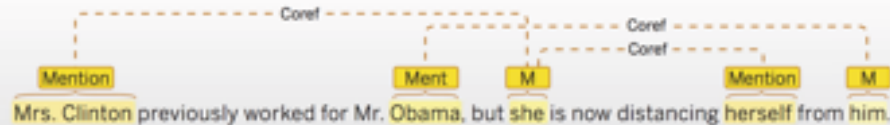
## Part of speech:




## Named entity recognition:



## Co-reference:




## Basic dependencies:



Advances in natural language processing,  
Hirschberg and Manning, Science 2015

# IBM Watson (question answering)

## machine translation



[https://www.youtube.com/watch?v=WFR3lOm\\_xhE](https://www.youtube.com/watch?v=WFR3lOm_xhE)



uncovering racial disparities in analyzing language  
from police body camera's (PNAS 2017)

# What makes language understanding hard?

## Polysemy

Today, I went to the **bank** to deposit a check.

The hut is located near the **bank** of the river.

Ardougne North **Bank** is near the **bank** of the River Dougne in the north part of the city.

Words can have  
multiple meanings



# What makes language understanding hard?

## Syntactic ambiguity

“One morning, I shot an elephant in my pajamas.”

“How he got in my pajamas, I don't know.”

(Groucho Marx)

**Who wore the pajamas?**



# What makes language understanding hard?

i know right	shake my head			for	your		
ikr	smh	he	asked	fir	yo	last	name
inter- jection	acronym	pronoun	verb	preposition	determiner	adjective	noun
				you		facebook	
so	he	can	add	u	on	fb	lololol
preposition	pronoun	verb	verb	pronoun	preposition	proper noun	inter- jection

# What makes language understanding hard?

*A string may have many possible interpretations in different contexts, and resolving ambiguity correctly may rely on knowing a lot about the world. (Noah Smith)*

- Linguistic **diversity**: languages, dialects, registers, styles
- Language is constantly **changing**

**ML for NLP:** What assumptions should the machine learning model have? (*How does language work?*)  
How should we represent language data in our models?

# Today: Vector representations of words and documents

**Represent** *documents, words, phrases, sentences, etc..* as vectors



→ [0.20, 0.80, 0.10, ..., -0.10]

dog → [0.10, 0.90, -0.20, ..., -0.40]

- What are the dimensions of the vector space?  
*One dimension for each word.*
- How do we measure distances in the vector space?

*To find similar documents, find similar words, as input to ML models, ...*

# Tokenization

**A tokenizer segments text into a sequence of tokens**

A  
tokenizer  
segments  
text  
into  
a  
sequence  
of  
tokens

We often take words as the basic level  
of analysis

# Tokenization: tokens vs types

**A tokenizer segments text into a sequence of tokens**

**Token:**

A  
tokenizer  
segments  
text  
into  
a  
sequence  
of  
tokens

**Type:**

a  
tokenizer  
segments  
text  
into  
sequence  
of  
tokens

# Tokenization challenges

始めまして。お元気ですか。

Japanese: No spaces between words

**Note:** Many libraries provide tokenization methods (e.g. nltk, spacy, scikit-learn)

*New York, European Union*

Multiword expressions

*I'm, don't*

Contractions

# Pre-processing steps

## Stop word removal

- *a, an, the, it, ..*
- Using a stop word list or by filtering words that appear in many documents

## Lemmatization

- *sang, sung, sings → sing*

## Stemming (strip endings of the word)

- E.g., *running* to *run*

## Lowercasing

- E.g., *Running* to *running*

## Removing infrequent words

- E.g., occurring in less than 10 (or 25, or ...) documents

**Overall goal:** reduce the number of *types* (usually thousands to millions in a large dataset!)

But could remove useful signals!

*Example: function words (the, a, and, however, on, etc..) are not so important for topic classification, but strong signals for authorship identification.*



# Edit distance

The minimum **edit distance** between two strings:

the minimum number of editing operations (insertion, deletion, substitution) needed to transform one string into another

INTENTION  
| | | | | | | | |  
\*EXECUTION  
d s s   i s

d = deletion  
s = substitution  
i = insertion

**Levenshtein distance:**

each of these operators has cost 1

# Bag of words representation

Only look at the presence (or frequency) of words, i.e. ignore the order. (*Often a surprisingly hard baseline to beat!*)

this is a  
nice restaurant



But..

Dog bites man



Man bites dog

**Question:** Come up with a task for which bag of words is probably sufficient. And another task for which it is not.

Sufficient: topic classification

Not sufficient: machine translation

# Jaccard similarity

## Jaccard similarity

Simple



$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

this is not bad  
this is really bad

Common (3): this, is, bad

Union (5): this, is, not, bad, really

Jaccard similarity: 3/5

Infrequent words (e.g. *bad*) versus frequent words (e.g. *is*)



Frequency and order of words are ignored

# Vector representations of documents

**word-document  
matrix**

	doc <sub>1</sub>	doc <sub>2</sub>	doc <sub>3</sub>	doc <sub>4</sub>	doc <sub>5</sub>	doc <sub>6</sub>	doc <sub>7</sub>
cat	5	2	0	1	4	0	0
dog	7	3	1	0	2	0	0
car	0	0	1	3	2	1	1

vector

Euclidian distance

$$\sqrt{\sum (a_i - b_i)^2}$$

**doc<sub>1</sub> and doc<sub>2</sub>**


Jaccard similarity = 1

Euclidian distance = 5

**doc<sub>1</sub> and doc<sub>5</sub>**

Jaccard similarity = 2/3

Euclidian distance = 5.477




RECAP!

# Euclidian distance

Distance between endpoint of vectors

**Euclidean distance:**

$$\|a - b\|_2 = \sqrt{\sum (a_i - b_i)^2}$$



**a** is closer (more similar) to **c** than to **b**, is this what we want?

**Extreme example:**

Concatenate two documents.

# Dot product

**Dot product:**

$$\mathbf{a} \cdot \mathbf{b} = \sum a_i b_i$$

$$\begin{array}{ll} \mathbf{a} = [0, 3, 5] & \mathbf{a} \cdot \mathbf{b} = 0 \times 3 + 3 \times 1 \\ \mathbf{b} = [3, 1, 2] & \quad \quad \quad + 5 \times 2 = 13 \end{array}$$

$$\begin{array}{ll} \mathbf{c} = [0, 1, 0] & \mathbf{c} \cdot \mathbf{d} = 0 \times 1 + 1 \times 1 \\ \mathbf{d} = [1, 1, 1] & \quad \quad \quad + 0 \times 1 = 1 \end{array}$$

If we have a **set-based** representation, i.e.  $\mathbf{a}$  and  $\mathbf{b}$  are **binary** vectors (1 = word is present, 0=absence), then the dot product is the number of words present in both documents (intersection)

*Compare to Jaccard Similarity!  
(normalization by dividing by  
the union of words)*

# Dot product

**Dot product:**

$$\mathbf{a} \cdot \mathbf{b} = \sum a_i b_i$$

**Length of a vector:**

**RECAP!**

$$\|\mathbf{a}\|_2 = \sqrt{\sum a_i^2}$$

Therefore:

$$\|\mathbf{a}\|_2 = \sqrt{\mathbf{a} \cdot \mathbf{a}}$$

**Normalization of a vector  
to unit length**

$$\frac{\mathbf{a}}{\|\mathbf{a}\|_2}$$

$$\mathbf{a} = [0, 3, 5]$$


$$\|\mathbf{a}\|_2 = \text{sqrt}(9 + 25) = 5.831$$

$$\frac{\mathbf{a}}{\|\mathbf{a}\|_2} = [0, 0.5145, 0.8575]$$

# Cosine similarity

Cosine similarity

$$\cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|} = \frac{\sum a_i b_i}{\sqrt{\sum a_i^2} \sqrt{\sum b_i^2}}$$



## Question:

What is the cosine distance between **a** & **b**

**a** = [0, 3, 5]

**b** = [3, 1, 2]

$$\mathbf{a} \cdot \mathbf{b} = 13$$

$$\|\mathbf{a}\| = \text{sqrt}(0^2 + 3^2 + 5^2) = \text{sqrt}(34)$$

$$\|\mathbf{b}\| = \text{sqrt}(3^2 + 1^2 + 2^2) = \text{sqrt}(14)$$

$$\rightarrow 13 / (\text{sqrt}(34) * \text{sqrt}(14)) \approx 0.60$$



# Cosine similarity

Cosine similarity


$$\cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|} = \frac{\sum a_i b_i}{\sqrt{\sum a_i^2} \sqrt{\sum b_i^2}}$$

When  $\mathbf{a}$  and  $\mathbf{b}$  are normalized:

$$\cos(\theta) = \mathbf{a} \cdot \mathbf{b}$$

If the vectors are orthogonal:

$$\mathbf{a} \cdot \mathbf{b} = 0$$




Cosine ranges from -1 (vectors pointing in opposite directions) to 0 (orthogonal) to 1 (vectors pointing in the same direction).

Raw frequencies (non-negative): 0-1

# Comparisons

## Cosine similarity

$$\cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|} = \frac{\sum a_i b_i}{\sqrt{\sum a_i^2} \sqrt{\sum b_i^2}}$$



## Euclidean distance

$$\|\mathbf{a} - \mathbf{b}\|_2 = \sqrt{\sum (a_i - b_i)^2}$$

If **a** and **b** have unit length, then Euclidean distance and cosine similarity will result in the same ordering (but reversed)

# Today: Vector representations of words and documents

**Represent** *documents, words, phrases, sentences, etc..* as vectors



→ [0.20, 0.80, 0.10, ..., -0.10]

dog → [0.10, 0.90, -0.20, ..., -0.40]

## Document representations

- What are the dimensions of the vector space?


One dimension for each word.  
Values: binary (presence or absence), frequency, weighting schemes (e.g. tf-idf, not discussed)

- How do we measure distances in the vector space?

Cosine similarity

# Today: Vector representations of words and documents

**Represent** *documents, words, phrases, sentences, etc..* as vectors

 → [0.20, 0.80, 0.10, ..., -0.10]

dog → [0.10, 0.90, -0.20, ..., -0.40]


dog and puppy,  
Monday and Tuesday,  
buy and purchase



Word overlap is not enough!  
Can we also map  
words to vectors??

- What are the dimensions of the vector space?
- How do we measure distances in the vector space?

# Word representations



*The standard way to represent  
the meaning of words in (neural) NLP*

# Word representations



when your	<b>cat</b>	leaves a mangled mouse
to care for your	<b>cat</b>	or kitten including food
the domestic	<b>cat</b>	is a small, typically furry
choosing to get a	<b>cat</b>	or dog is a

*You shall know a word by  
the company it keeps*  
(Firth, J. R. 1957:11)

To assign similar vectors to similar words a notion of similarity is needed.

**The distributional hypothesis:** Words that occur in similar contexts tend to have similar meanings.

# Vector representations of words

		doc <sub>1</sub>	doc <sub>2</sub>	doc <sub>3</sub>	doc <sub>4</sub>	doc <sub>5</sub>	doc <sub>6</sub>	doc <sub>7</sub>	
<b>documents as context</b> <b>word-document</b> <b>matrix</b>	cat	5	2	0	1	4	0	0	} vector
	dog	7	3	1	0	2	0	0	
	car	0	0	1	3	2	1	1	

# Vector representations of words

		doc <sub>1</sub>	doc <sub>2</sub>	doc <sub>3</sub>	doc <sub>4</sub>	doc <sub>5</sub>	doc <sub>6</sub>	doc <sub>7</sub>
<b>documents as context</b> <b>word-document</b> <b>matrix</b>	cat	5	2	0	1	4	0	0
	dog	7	3	1	0	2	0	0
	car	0	0	1	3	2	1	1

neighboring words  
as context

word-word  
matrix

		cat	dog	car	bike	book	housetree	
	cat	0	3	1	1	1	2	3
	dog	3	0	2	1	1	3	1
	car	1	2	0	3	2	2	0



# Vector representations of words

## Properties:

- Vectors are sparse:  
Many zero entries
- Values are based on  
frequencies  
(sometimes  
weighted)

**documents as context**

**word-document  
matrix**

	doc <sub>1</sub>	doc <sub>2</sub>	doc <sub>3</sub>	doc <sub>4</sub>	doc <sub>5</sub>	doc <sub>6</sub>	doc <sub>7</sub>
cat	5	2	0	1	4	0	0
dog	7	3	1	0	2	0	0
car	0	0	1	3	2	1	1

**neighboring words  
as context**

**word-word  
matrix**

	cat	dog	car	bike	book	housetree	
cat	0	3	1	1	1	2	3
dog	3	0	2	1	1	3	1
car	1	2	0	3	2	2	0

# Word embeddings (Dense word vectors)

## Word embeddings:

- Vectors are dense
- The values are learned by **training a classifier** to distinguish nearby and far away words
- Individual dimensions are less interpretable.


Dense real-valued vectors

cat 

0.52	0.84	0.01	...	0.23
------	------	------	-----	------

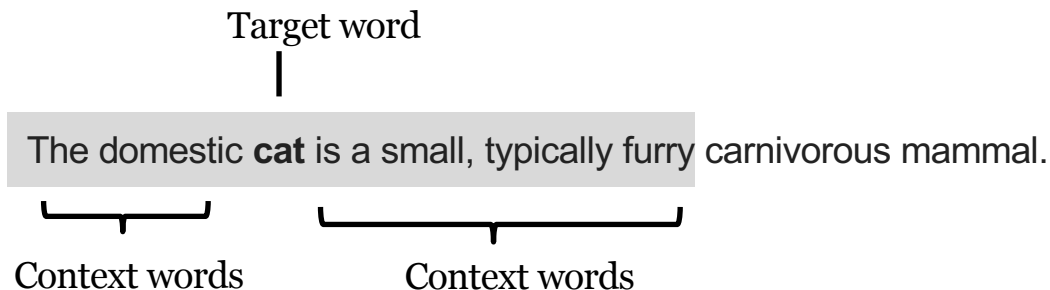
dog 

0.40	0.90	0.10	...	0.40
------	------	------	-----	------



# How are word embeddings learned?

(the skipgram model)



Embedding vectors  
are essentially a by-  
product!

target word	context word	label
cat	small	1
cat	furry	1
cat	car	0
...	...	....

## Different methods based on similar ideas:

- Word2vec (Mikolov et al. 2013), skipgram & continuous bag of words.  
(<https://code.google.com/archive/p/word2vec/>, another implementation in <https://radimrehurek.com/gensim/>)
- GloVe  
(<https://nlp.stanford.edu/projects/glove/>)
- FastText (<https://fasttext.cc/>)

# Properties of word embeddings

We can use cosine similarity to find similar words in the vector space.

san francisco



los_angeles	0.666175
golden_gate	0.571522
oakland	0.557521
california	0.554623



<https://code.google.com/archive/p/word2vec/>  
[https://en.wikipedia.org/wiki/San\\_Francisco](https://en.wikipedia.org/wiki/San_Francisco)

# Properties of word embeddings

We can look at analogies in the vector space



*king - man + woman  $\approx$  queen*

Linguistic Regularities in  
Continuous Space Word  
Representations, Mikolov et  
al. 2013

# Word embeddings (Dense word vectors)

Dense real-valued vectors


**cat**

0.52 0.84 0.01 .... 0.23

**dog**

0.40 0.90 0.10 .... 0.40

How are these word representations used?




*In neural networks  
(text classification,  
sequence tagging,  
etc..)*

*As research  
objects*



# Using embeddings in ML models



**Bag of words model** (surprisingly effective). Average over the vectors! (resulting in a vector of size 7)

Most commonly the word vectors are used as input in **neural networks**

# Word embeddings (Dense word vectors)

Dense real-valued vectors


**cat**

0.52 0.84 0.01 .... 0.23

**dog**

0.40 0.90 0.10 .... 0.40

How are these word representations used?




*In neural networks  
(text classification,  
sequence tagging,  
etc..)*

*As research  
objects*






# Measuring change using embeddings




pre-1970



post-1970

# Empirical data: glo

August 2013 rapper Chief Keef released “Gotta Glo Up One Day.”



[Shoemark, Liza, et al. EMNLP 2019]

# *What do you need to know*

- K-nearest neighbor method (algorithm, pros and cons, effect of K, distance measures)
- You should be able to compute Manhattan/L1 and Euclidian/L2 distance, cosine similarity, dot product
- The frequently used preprocessing steps in natural language processing
- What makes processing and analyzing language difficult?
- Representing documents (from documents to vectors)
- Representing words (from words to vectors)

# NLP tools

- **Natural Language Toolkit (NLTK)**

<http://www.nltk.org/>



- **Spacy** <https://spacy.io/>

- Many machine learning libraries (e.g. scikit-learn) also implement basic NLP methods.

**spaCy**

# Books (drafts online!)

- Speech & language processing (3<sup>rd</sup> edition, to appear in 2019) by Jurafsky and Martin
  - <https://web.stanford.edu/~jurafsky/slp3/>
- Introduction to Natural Language Processing (1<sup>st</sup> edition, to appear in 2019) by Jacob Eisenstein
  - <https://github.com/jacobeisenstein/gt-nlp-class/blob/master/notes/eisenstein-nlp-notes.pdf>

# Thanks

Some slides based on (or inspired by) slides by  
Matt Gormley, Carlos Guestrin, Soheil Feizi,  
Victor Lavrenko, Marine Carpuat, Noah Smith,  
Marijn Schraagen